

CIS341

Floating point representation

Bryan S. Kim

Assistant Professor

Syracuse University

College of Engineering & Computer Science
Department of Electrical Engineering & Computer Science

Sep. 14, 2023

Slides based on contents from CS61C @ UC Berkeley
& 15-213 @ CMU

Last lecture: integer representation

Crux of the problem: we can't have infinite bits

Two's complement representation

- Limited number of bit, potential overflow

Signed vs. unsigned

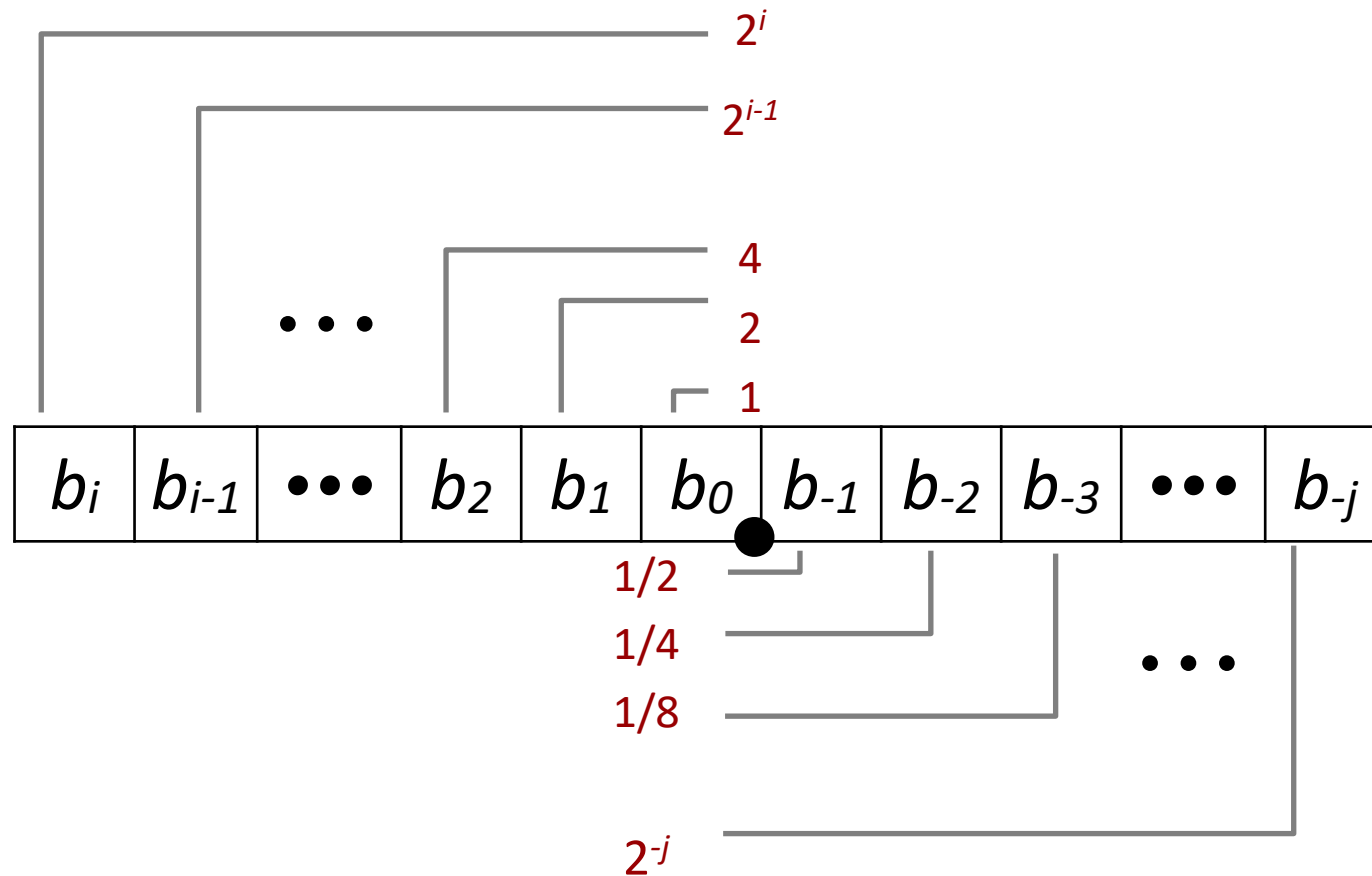
- Same bit representation
- Different interpretation
- Sometimes different operations (example: right shift)
- What happens when they are mixed

Fractional binary numbers

What is 1011.101_2 ?

Fractional binary numbers

Bits right of the binary point represent fractional powers of 2



Fractional binary numbers: example

Value	Representation	
$5 \text{ \& } 3/4 = 23/4$	101.11_2	$= 4 + 1 + 1/2 + 1/4$
$2 \text{ \& } 7/8 = 23/8$	10.111_2	$= 2 + 1/2 + 1/4 + 1/8$
$1 \text{ \& } 7/16 = 23/16$	1.0111_2	$= 1 + 1/4 + 1/8 + 1/16$

Observations

- Divide by 2 by shifting right
- Multiply by 2 by shifting left
- Numbers of form $0.11111\dots_2$ are almost 1.0
 - $1/2 + 1/4 + 1/8 + \dots + 1/2^i$ converge to 1.0

Representable numbers

Limitation #1

- Can only exactly represent numbers of the form $x/2^k$
 - In decimal, $1/3$ is $0.3333333...$
 - In binary, $1/3$ is $0.010101010101...$


Limitation #2

- Just one setting of binary point within the w -bits

The importance of floating point

Missile interception failure on Feb. 25, 1991

- Internal system clock incremented every 0.1 seconds, and to compute 1 second, counter was multiplied by 1/10.
- $1/10 = 0.0001100110011001100..._2$, and this was stored as 24-bits.
- Introduced an error of $0.000000000000000000000000011001100_2$, or 0.0000000095 seconds for every 0.1 second.
- After 100 hours of operation, the error was large enough to miss the interception.



IEEE floating point (IEEE 754)

Established in 1985 as a uniform standard for floating point

Supported by all major CPUs

- Some don't fully implement IEEE 754 (some GPUs)

Driven by numerical concerns

- Numerical analysts predominated over hardware designers in defining the standard
- Hard to make fast in hardware

Floating point representation

$$341_{10} = (-1)^0 \times 1.01010101 \times 2^8$$

Numerical form

- $(-1)^s \times M \times 2^E$
- Sign bit **s** determines whether the number is negative or positive
- Mantissa **M** is normally a fractional value in range of [1.0, 2.0)
- Exponent **E** weighs value by power of two

Encoding

- MSB **s** is sign bit **s**
- **exp** field encodes **E** (but is not equal to E)
- **frac** field encodes **M** (but is not equal to M)



Precision options

Single precision (32-bits): float

≈ 7 decimal digits, $10^{\pm 38}$



Double precision (64-bits): double

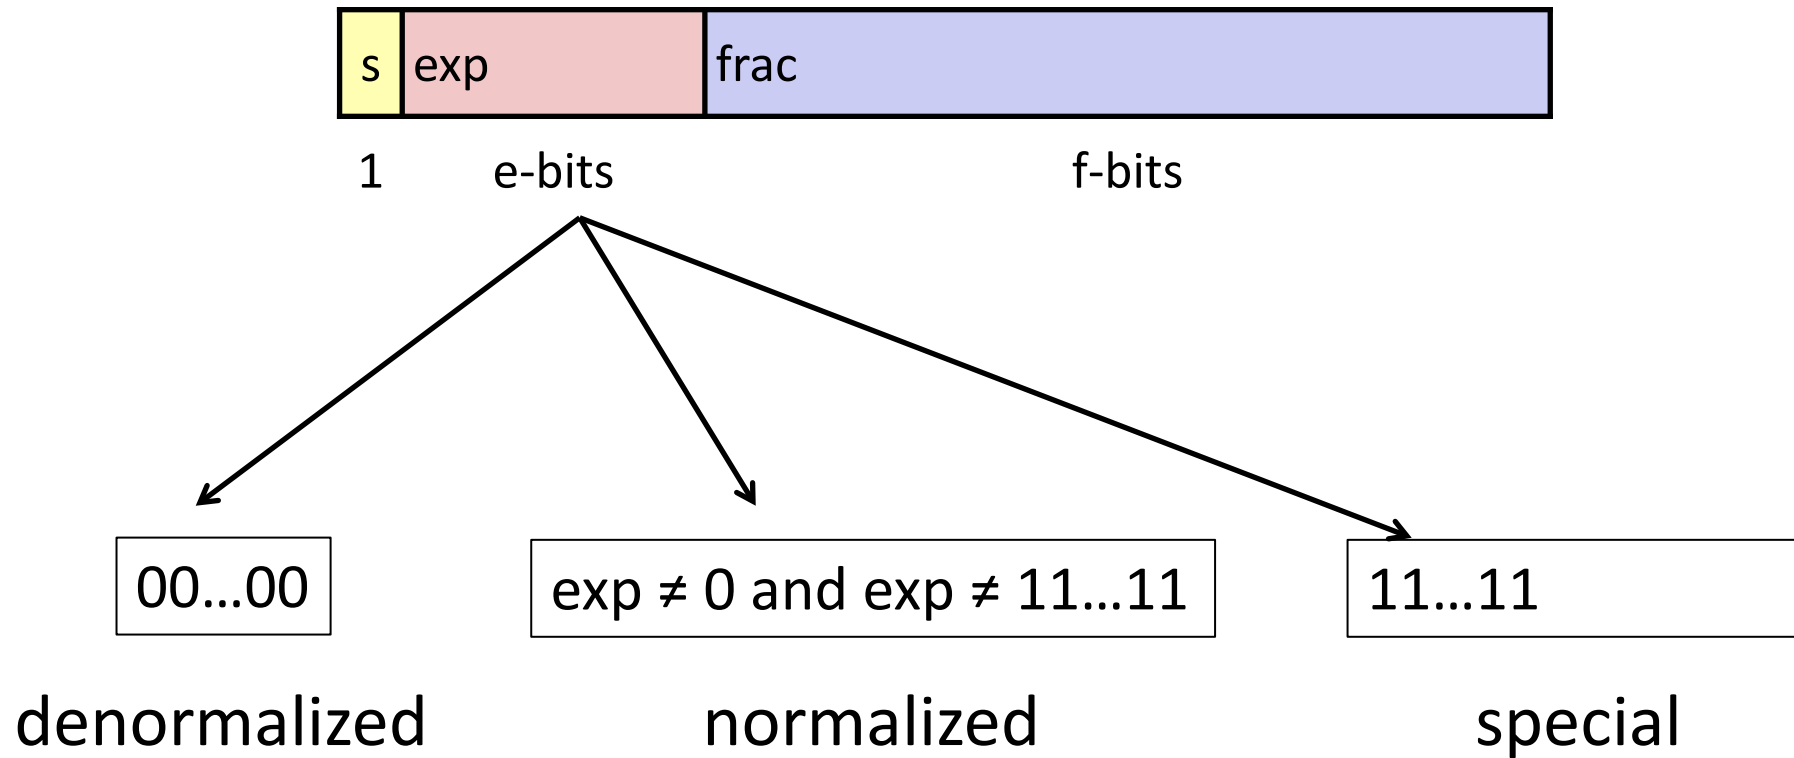
≈ 16 decimal digits, $10^{\pm 308}$




Other formats exist

- Half precision (16-bits)
- Quad precision (128-bits): long double

Three kinds of floating point numbers



Visualization of the floating point encodings



Normalized values

$$v = (-1)^s M 2^E$$

When $\text{exp} \neq 000\dots 0$ and $\text{exp} \neq 111\dots 1$

Exponent coded as a biased value: $E = \text{exp} - \text{Bias}$

- **exp** : unsigned value of **exp** field
- Bias: $2^{k-1} - 1$, where k is the number of exponent bits
 - Single precision: 127 (**exp** : 1 to 254, E : -126 to 127)
 - Double precision: 1023 (**exp**: 1 to 2046, E : -1022 to 1023)

Mantissa coded with implied leading 1: $M = 1.\text{xxx}\dots\text{x}_2$

- xxx.x: bits of the **frac** field
- Minimum when **frac** is 000...0 ($M = 1.0$)
- Maximum when **frac** is 111...1 ($M \approx 2.0$)

$$v = (-1)^s M 2^E$$

$$E = \text{exp} - \text{Bias}$$

Normalized encoding example

`float f = 341.0;`

$$341 = 101010101_2 = 1.01010101 \times 2^8$$


Mantissa

- M = 1.01010101
- frac = 010101010000000000000000₂

Exponent

- E = 8
- Bias = 127
- Exp = 135 = 10000111₂

Result



Denormalized values

$$v = (-1)^s M 2^E$$
$$E = 1 - \text{Bias}$$

When $\text{exp} = 000\dots 0$

Exponent value: $E = 1 - \text{Bias}$

- Bias is same as the case for normalized ($2^{k-1}-1$)

Mantissa coded with implied leading 0: $M = 0.\text{xxx}\dots\text{x}_2$

- xxx.x : bits of the **frac** field

Examples

- **exp** = $000\dots 0$, **frac** = $000\dots 0$
 - Represents zero value (+0 when $s = 0$, -0 when $s=1$)
- **exp** = $000\dots 0$, **frac** $\neq 000\dots 0$
 - Equi-spaced numbers close to zero

Special values

When $exp = 111...1$

When **exp** = 111...1, **frac** = 000...0

- Represents infinity (∞)
- + ∞ when $s = 0$, - ∞ when $s = 1$

When **exp** = 111...1, **frac** \neq 000...0

- Not-a-number (NaN)
- Example: $\sqrt{-1}$, $\infty - \infty$, $\infty \times 0$

C float decoding example

Float: 0xC0A00000

C float decoding example


Float: 0xC0A00000

Binary: 1100 0000 1010 0000 0000 0000 0000 0000

C float decoding example

Float: 0xC0A00000

Binary: 1100 0000 1010 0000 0000 0000 0000 0000



exp is neither 000...0 or 111...1

Normalized value


$E = \text{exp} - \text{Bias}$

$M = 1.\text{xxx}$

C float decoding example

Float: 0xC0A00000

Binary: 1100 0000 1010 0000 0000 0000 0000 0000



E = exp - Bias = 129 - 127 = 2

S = 1

M = 1.010 0000 0000 0000 0000 0000
= 1 + 1/4 = 1.25

exp is neither 000...0 or 111...1

Normalized value


E = exp - Bias

M = 1.xxx

C float decoding example

Float: 0xC0A00000

Binary: 1100 0000 1010 0000 0000 0000 0000 0000



$$E = \text{exp} - \text{Bias} = 129 - 127 = 2$$

$$S = 1$$

$$M = 1.010\ 0000\ 0000\ 0000\ 0000\ 0000$$

$$= 1 + 1/4 = 1.25$$

$$v = (-1)^S \times M \times 2^E = (-1)^1 \times 1.25 \times 2^2 = -5$$

exp is neither 000...0 or 111...1

Normalized value

$E = \text{exp} - \text{Bias}$

$M = 1.\text{xxx}$

C float decoding example

Float: 0x001C0000

C float decoding example


Float: 0x001C0000

Binary: 0000 0000 0001 1100 0000 0000 0000 0000

C float decoding example

Float: 0x001C0000

Binary: 0000 0000 0001 1100 0000 0000 0000 0000




exp is 000...0
Denormalized value
 $E = 1 - \text{Bias}$
 $M = 0.\text{xxx}$

C float decoding example

Float: 0x001C0000

Binary: 0000 0000 0001 1100 0000 0000 0000 0000



E = 1 - Bias = 1 - 127 = -126

S = 0

M = 0.001 1100 0000 0000 0000 0000

= $1/8 + 1/16 + 1/32 = 7/32$

exp is 000...0

Denormalized value


$E = 1 - \text{Bias}$

$M = 0.\text{xxx}$

C float decoding example

Float: 0x001C0000

Binary: 0000 0000 0001 1100 0000 0000 0000 0000



$$E = 1 - \text{Bias} = 1 - 127 = -126$$

$$S = 0$$

$$M = 0.001\ 1100\ 0000\ 0000\ 0000\ 0000$$

$$= 1/8 + 1/16 + 1/32 = 7/32$$

$$v = (-1)^S \times M \times 2^E = (-1)^0 \times 7/32 \times 2^{-126} \approx 2.571393892 \times 10^{-39}$$

exp is 000...0

Denormalized value

$E = 1 - \text{Bias}$

$M = 0.\text{xxx}$

Group activity (1 of 2)

Floating point arithmetic

Floating point operations: basic idea

$$\mathbf{x} +_{\mathbf{f}} \mathbf{y} = \text{Round}(\mathbf{x} + \mathbf{y})$$

$$\mathbf{x} \times_{\mathbf{f}} \mathbf{y} = \text{Round}(\mathbf{x} \times \mathbf{y})$$

First **compute the exact result**

Make it fit into desired precision

- Possibly ends up being infinity if exponent too large
- Possibly round to fit into frac

Rounding

Should 0.5 round to 0 or 1?

	1.4	1.6	1.5	2.5	-1.5
Towards zero	1 ↓	1 ↓	1 ↓	2 ↓	-1 ↑
Round down	1 ↓	1 ↓	1 ↓	2 ↓	-2 ↓
Round up	2 ↑	2 ↑	2 ↑	3 ↑	-1 ↑
Nearest even*	1 ↓	2 ↑	2 ↑	2 ↓	-2 ↓

*Round to nearest, but if half-way in-between then round to nearest even

Round-to-even

Default rounding mode (other roundings are biased)

If exactly halfway between two possible values

- Round so that the least significant digit is even

Otherwise, round to nearest

Decimal examples (round to nearest hundredth)

7.8749999	7.87	(less than halfway)
7.8750000	7.88	(half way: round up to even)
7.8750001	7.88	(greater than halfway)
7.8800000	7.88	
7.8849999	7.88	(less than halfway)
7.8850000	7.88	(half way: round down to even)

Rounding binary

“Even” when least significant bit is 0

“Half way” when bits right of the rounding position = 100...0

Binary example (round to nearest 1/4)

10.00011	10.00	(less than halfway)
10.00100	10.00	(halfway: round down)
10.00101	10.01	(more than halfway)
10.01000	10.01	
10.01011	10.01	(less than halfway)
10.01100	10.10	(halfway: round up)

Floating point multiplication

$$(-1)^{S1} M1 2^{E1} \times (-1)^{S2} M2 2^{E2}$$

Result format: $(-1)^S M 2^E$

- Sign $s = s1 \wedge s2$ (note: $(-1)^{s1 \wedge s2}$ is same as $(-1)^{s1+s2}$)
- Mantissa $M = M1 \times M2$
- Exponent $E = E1 + E2$

Adjustments

- If $M \geq 2$, shift M right, increment E
- If E is out of range, overflow into infinity
- Round M to fit frac precision

Example (given 4-bit mantissa)

$$\begin{aligned} 1.010 \times 2^2 \times 1.110 \times 2^3 &= 10.0011 \times 2^5 \\ &= 1.00011 \times 2^6 = 1.001 \times 2^6 \end{aligned}$$

Floating point addition

$$(-1)^{S1} M1 2^{E1} + (-1)^{S2} M2 2^{E2} \quad (\text{assume } E1 > E2)$$

Result format: $(-1)^S M 2^E$

- Align and add to get sign s and mantissa M
- Exponent E is $E1$ (the larger of the two)


Adjustments

- If $M \geq 2$, shift M right, increment E
- If $M < 1$, shift M left, decrement E
- If E is out of range, overflow into infinity
- Round M to fit frac precision

Example (given 4-bit mantissa)

$$\begin{aligned} 1.010 \times 2^2 + 1.110 \times 2^3 &= (0.101 + 1.110) \times 2^3 \\ &= 10.011 \times 2^3 = 1.010 \times 2^4 \end{aligned}$$

Get binary points lined up



Floating point arithmetic quirks

$$(3.14 + 1E10) - 1E10 = ?$$

$$3.14 + (1E10 - 1E10) = ?$$

$$1E20 * (1E20 - 1E20) = ?$$

$$1E20 * 1E20 - 1E20 * 1E20 = ?$$

Floating point in C

`float` *single precision (32-bit)*

`double` *double precision (64-bit)*

Casting between `int`, `float`, and `double` changes bit representation

`double` or `float` into `int`

- Truncates fractional part

`int` into `double`

- Exact conversion (`int` is 32-bits, `double`'s mantissa is 52-bits)

`int` into `float`

- Will round-to-even

Floating point casting quirks

Given `int x; float f; double d;`

False	<code>x == (int) (float) x</code>	
True	<code>x == (int) (double) x</code>	
True	<code>f == (float) (double) f</code>	
False	<code>d == (double) (float) d</code>	
True	<code>f == -(-f)</code>	
False	<code>2/3 == 2/3.0</code>	
True	<code>d < 0.0</code>	$\rightarrow ((d*2) < 0.0)$
True	<code>d > f</code>	$\rightarrow -f > -d$
True	<code>d * d >= 0.0</code>	
False	<code>(d+f)-d == f</code>	

Group activity (2 of 2)

To-dos

proj0 (due Yesterday!)

hw0 (due next Wednesday)